FINOVA CREDIT RISK PREDICTION USING DATAMINING

A PROJECT REPORT SUBMITTED IN THE PARTIAL FULFILLMENT OF THE REQUIREMENTS FOR THE AWARD OF THE DEGREE OF

BACHELOR OF TECHNOLOGY

IN

DEPARTMENT OF COMPUTER SCIENCE ENGINNERING

BY

NAME OF THE STUDENT	ROLL NUMBER
KUNAL KUMAR DAPPU	2320030167

UNDER THE ESTEEMED GUIDANCE OF

DR P.KRISHNA KISHORE
ASSISTANT PROFESSOR

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

Koneru Lakshmaiah Education Foundation
(Deemed to be University estd. u/s. 3 of the UGC Act, 1956)
Off-Campus: Bachupally-Gandimaisamma Road, Bowrampet, Hyderabad, Telangana - 500 043.
Phone No: 7815926816, www.klh.edu.in

K L (DEEMED TO BE) UNIVERSITY
DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

&

DEPARTMENT OF COMPUTER SCIENCE AND INFORMATION TECHNOL

DECLARATION

THE PROJECT REPORT ENTITLED " **FINOVA CREDIT RISK PREDICTION USING DATAMINING**" IS A RECORD OF BONAFIDE WORK OF **KUNAL KUMAR DAPPU – 2320030167,** SUBMITTED IN PARTIAL FULFILLMENT FOR THE AWARD OF B. TECH IN COMPUTER SCIENCE AND ENGINEERING (OR) COMPUTER SCIENCE AND INFORMATION TECHNOLOGY TO THE K L UNIVERSITY. THE RESULTS EMBODIED IN THIS REPORT HAVE NOT BEEN COPIED FROM ANY OTHER DEPARTMENTS/UNIVERSITY/INSTITUTE.

KUNAL KUMAR DAPPU – 23200300167

K L (DEEMED TO BE) UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

CERTIFICATE

THIS IS CERTIFY THAT THE MINI PROJECT BASED REPORT ENTITLED " **FINOVA CREDIT RISK PREDICTION USING DATAMINING**" IS A BONAFIDE WORK DONE AND SUBMITTED BY **KUNAL KUMAR DAPPU – 2320030167** IN PARTIAL FULFILLMENT OF THE REQUIREMENTS FOR THE AWARD OF THE DEGREE OF **BACHELOR OF TECHNOLOGY** IN DEPARTMENT OF COMPUTER SCIENCE ENGINEERING, K L (DEEMED TO BE UNIVERSITY), DURING THE ACADEMIC YEAR **2025-2026.**

SIGNATURE OF THE GUIDE

SIGNATURE OF THE COURSE COORDINATOR

SIGNATURE OF THE HOD

ACKNOWLEDGEMENT

THE SUCCESS IN THIS PROJECT WOULD NOT HAVE BEEN POSSIBLE BUT FOR THE TIMELY HELP AND GUIDANCE RENDERED BY MANY PEOPLE. OUR WISH TO EXPRESS MY SINCERE THANKS TO ALL THOSE WHO HAS ASSISTED US IN ONE WAY OR THE OTHER FOR THE COMPLETION OF MY PROJECT.

OUR GREATEST APPRECIATION TO MY COURSE COORDINATOR **DR P.KRISHNA KISHORE,** DEPARTMENT OF COMPUTER SCIENCE WHICH CANNOT BE EXPRESSED IN WORDS FOR HIS/HER TREMENDOUS SUPPORT, ENCOURAGEMENT AND GUIDANCE FOR THIS PROJECT.

WE EXPRESS OUR GRATITUDE TO **DR. P. VENKATESHWARA RAO (CSE)** HEAD OF THE **DEPARTMENT FOR COMPUTER SCIENCE ENGINEERING** FOR PROVIDING US WITH ADEQUATE FACILITIES, WAYS AND MEANS BY WHICH WE ARE ABLE TO COMPLETE THIS PROJECT-BASED LAB.

WE THANK ALL THE MEMBERS OF TEACHING AND NON-TEACHING STAFF MEMBERS, AND ALSO WHO HAVE ASSISTED ME DIRECTLY OR INDIRECTLY FOR SUCCESSFUL COMPLETION OF THIS PROJECT. FINALLY, I SINCERELY THANK MY PARENTS, FRIENDS AND CLASSMATES FOR THEIR KIND HELP AND COOPERATION DURING MY WORK.

KUNAL KUMAR DAPPU – 23200300167

ABSTRACT

The financial services industry is increasingly dependent on digital lending platforms such as NBFCs, BNPL systems, and fintech applications, all of which require fast and reliable borrower risk assessment. The attached presentation frames the project as a credit risk prediction system for Indian FinTech lending apps and emphasizes the need for accurate, explainable, and deployment-ready risk scoring for digital-first lending ecosystems.

FINOVA — AI-Driven Credit Risk Engine is designed as an end-to-end analytical system that integrates synthetic credit data generation, preprocessing, feature engineering, machine learning, risk banding, decision rules, explainability, and interactive dashboard presentation. The presentation identifies core components including logistic regression, XGBoost or LightGBM, risk band assignment, decision output, and SHAP-based explanation, and this report expands those components into a formal project document.

The proposed system addresses the limitations of traditional credit evaluation by reducing manual effort, improving scalability, and supporting consistent model-based decisions. The workflow begins with applicant and loan data, transforms those inputs into engineered features, predicts the probability of default, maps the prediction into business-oriented risk bands, generates an approve, review, or reject decision, and presents the outcome through an interactive dashboard.

The report is organized using the same broad chapter style as the uploaded academic sample, which includes institutional front matter followed by introduction, requirements, architecture, methodology, implementation, results, limitations, future enhancements, conclusion, and appendices. This structure makes the document suitable for expansion into a 50-page academic report by adding screenshots, diagrams, tables, model outputs, and implementation evidence.

TABLE OF CONTENTS

1. Introduction
   - 1.1 Background
   - 1.2 Problem Statement
   - 1.3 Objective
   - 1.4 Scope
   - 1.5 Importance of the Project
2. System Overview
   - 2.1 Project Name
   - 2.2 Key Components
   - 2.3 Workflow
3. System Architecture
   - 3.1 Architecture Overview
   - 3.2 Architecture Description
   - 3.3 Layered Flow
   - 3.4 Architecture Diagram
4. Data Warehouse Design
   - 4.1 Star Schema
   - 4.2 Fact Table
   - 4.3 Dimension Tables
   - 4.4 Advantages of Warehouse Design
5. Dataset Description
   - 5.1 Data Attributes
   - 5.2 Data Size
   - 5.3 Data Generation
   - 5.4 Feature Groups
   - 5.5 Data Dictionary Notes
6. System Requirements
   - 6.1 Hardware Requirements
   - 6.2 Software Requirements

7. Literature Survey
    7.1 Traditional Credit Scoring
    7.2 ML-Based Credit Risk Models
    7.3 Thin-File Borrower Challenges
    7.4 Explainability in Lending
    7.5 Research Gap and Novelty
8. Technology Stack
    8.1 Python
    8.2 Data Analysis Libraries
    8.3 Machine Learning Libraries
    8.4 Explainability Tools
    8.5 Streamlit
    8.6 Integration of Technologies
9. Methodology
    9.1 Data Cleaning
    9.2 Encoding
    9.3 Feature Engineering
    9.4 Train-Test Strategy
    9.5 Imbalance Handling
    9.6 Baseline Model
    9.7 Advanced Model
    9.8 Explainability Workflow
    9.9 End-to-End Pipeline
10. Machine Learning Models
    10.1 Logistic Regression
    10.2 XGBoost
    10.3 LightGBM
    10.4 Model Comparison Logic
11. Risk Banding and Decision Engine
    11.1 PD Threshold Design
    11.2 Low, Medium, High Risk Mapping

11.3 Approval Logic
11.4 Business Interpretation
12. Explainability (SHAP)
12.1 Purpose
12.2 Global Explanations
12.3 Individual Reason Codes
12.4 Compliance Relevance
13. Dashboard Design
13.1 Streamlit Interface
13.2 Input Module
13.3 Prediction Output Module
13.4 Visualization and Explanation Panel
13.5 Report Export Options
14. Implementation
14.1 Environment Setup
14.2 Data Preparation
14.3 Model Training
14.4 Threshold Tuning
14.5 Dashboard Integration
14.6 Testing and Validation
15. Model Evaluation
15.1 Evaluation Metrics
15.2 Why PR-AUC and KS Matter
15.3 Expected Results
15.4 Real-World Lending Interpretation
16. Applications
17. Advantages
18. Limitations
19. Future Enhancements
20. Conclusion
21. References

1. INTRODUCTION

The uploaded sample project report begins its technical body with an introduction chapter followed by structured subsections such as background, problem statement, scope, objectives, and application importance, and this report adopts the same academic sequencing. The attached presentation positions the project within India's digital lending ecosystem and highlights the need for accurate, explainable, and practical credit scoring for FinTech applications.

Digital lending has transformed the way borrowers access short-term and medium-term credit. Traditional manual underwriting is often too slow for instant lending products, while newer fintech and NBFC platforms must evaluate large numbers of applicants quickly and consistently. This shift creates a strong need for intelligent risk engines that can combine borrower attributes, repayment indicators, and digital behavior into timely lending decisions.

FINOVA is designed to address this requirement through an integrated credit risk prediction workflow. The system combines data engineering, predictive modeling, business rule mapping, and explainable AI so that a raw applicant profile can be translated into a practical lending decision with transparent reasoning.

1.1 BACKGROUND

The attached project presentation states that the financial industry has undergone rapid digital transformation through NBFC lending platforms, BNPL systems, and fintech applications that rely on data-driven loan approval workflows. In such environments, credit risk assessment is no longer a periodic back-office process; it is a real-time operational requirement that directly influences approval speed, portfolio quality, and customer experience.

The presentation also notes that poor credit decisions can result in higher default rates, financial loss, and weaker portfolio performance, which makes predictive accuracy and operational consistency critical for lending institutions. Borrowers in digital ecosystems often

include thin-file or new-to-credit users, so the system must infer risk even when conventional bureau history is limited.

1.2 PROBLEM STATEMENT

The attached content identifies several weaknesses of traditional credit evaluation, including manual processing, time consumption, human bias, and limited scalability. The presentation further frames the project around the need for an automated system that predicts credit risk accurately, provides explainable decisions, supports real-time analysis, and integrates both data engineering and machine learning.

From a business perspective, the challenge is not only to classify borrowers as good or bad but also to convert model output into operational actions such as approve, review, or reject. A production-oriented risk engine must therefore balance accuracy, interpretability, and business usefulness rather than focusing on a single metric in isolation.

1.3 OBJECTIVE

The stated objective of the project is to develop a FinTech Credit Risk Prediction System that uses machine learning, data warehousing, and interactive visualization to support intelligent lending decisions. Building on that objective, this report expands the system into an end-to-end framework that includes data generation, preprocessing, feature engineering, classification, thresholding, explainability, and dashboard-based delivery.

1.4 SCOPE

The project scope described in the attached content includes data generation and preprocessing, feature engineering, machine learning models, risk classification, SHAP-based explainability, and an interactive Streamlit dashboard. This report preserves that scope and organizes it into a formal academic document similar to the uploaded sample report structure. +1

The current project is intended as a realistic academic simulation of a credit risk platform rather than a direct production banking system. It demonstrates business-aware underwriting

logic and explainable ML, but it does not claim live integration with regulated lender infrastructure unless such implementation evidence is separately provided.

1.5 IMPORTANCE OF THE PROJECT

The project is important because digital lending ecosystems need immediate and reliable borrower evaluation while still maintaining auditability and business consistency. The attached presentation emphasizes deployment readiness, explainability, and the ability to serve thin-file users in India's digital lending landscape, which makes the project relevant to modern fintech operations.

2. SYSTEM OVERVIEW

The attached project description names the system **Finova — AI-Driven Credit Risk Engine** and identifies its core components as a data warehouse, ML model layer, risk banding engine, decision engine, and Streamlit dashboard. This aligns well with the layered, chapter-driven explanatory style used in the sample report's overview and architecture chapters. +1

2.1 PROJECT NAME

The project name is **FINOVA — AI-Driven Credit Risk Engine**. The title reflects the system's combination of financial analytics, predictive intelligence, and operational decision support.

2.2 KEY COMPONENTS

According to the attached content, the major components are the data warehouse, machine learning models such as Logistic Regression and XGBoost, a risk banding engine, a decision engine, and a Streamlit dashboard. The presentation also mentions LightGBM as an advanced model option and SHAP-based explanation for user-level reason codes and global importance.

2.3 WORKFLOW

The project workflow described in the attached content follows a clear sequence: the user enters loan details, the data is processed and transformed, the ML model predicts probability of default, a risk band is assigned, a decision is generated, and an explanation is provided.

This stepwise operational flow provides the backbone for the system design and implementation sections of the report.

3. SYSTEM ARCHITECTURE

The presentation explicitly states that the project follows a layered architecture consisting of a Data Layer, Processing Layer, ML Layer, Decision Layer, and Presentation Layer. The sample PDF also dedicates a separate chapter to architecture and a separate subsection to the architecture diagram, so the same documentation style is used here. +1

3.1 ARCHITECTURE OVERVIEW

The system architecture is designed as a modular lending intelligence pipeline in which raw applicant and loan information is converted into a risk-oriented business decision. Each layer performs a distinct function so that the system remains scalable, interpretable, and easier to maintain.

3.2 ARCHITECTURE DESCRIPTION

The Data Layer contains the synthetic lending dataset described in the presentation, which includes more than 25,000 tabular records representing borrower and credit-related attributes. The Processing Layer performs cleaning, encoding, and feature generation before transferring structured inputs to the modeling components.

The ML Layer includes Logistic Regression as a transparent benchmark and XGBoost or LightGBM as advanced learners capable of capturing non-linear patterns. The Decision Layer converts the predicted probability of default into business-friendly risk bands and lending actions, while the Presentation Layer surfaces the final result through a Streamlit interface with visual and explanatory elements.

3.3 LAYERED FLOW

A concise representation of the layered flow is as follows:

Applicant Input → Data Validation → Preprocessing → Feature Engineering → Model Scoring → PD Estimation → Risk Band Assignment → Decision Output → Explanation Display

3.4 ARCHITECTURE DIAGRAM

The attached presentation says the architecture workflow diagram shows the complete end-to-end system flow, including data sources, model development, decision outputs such as Approve, Reject, or Manual Review, and monitoring loops. In the final formatted report, this subsection should include the image or recreated diagram from the presentation so the report visually matches the sample report's architecture section style.

4. DATA WAREHOUSE DESIGN

The user-provided content states that the system includes a data warehouse designed using a star schema with a loan transaction fact table and dimension tables for applicant details, credit history, device information, and time. A warehouse layer of this kind supports analytical consistency and aligns with production-inspired fintech system design.

4.1 STAR SCHEMA

The warehouse follows a star schema because this design enables simple joins, cleaner analytical logic, and efficient reporting. The central fact table records lending transactions or scored applications, while the surrounding dimension tables provide descriptive context used for slicing, filtering, and business interpretation.

4.2 FACT TABLE

The fact table stores measurable loan-related events, such as application amount, predicted risk score, repayment outcome label, decision result, and event time. This table serves as the analytical center of the warehouse because it connects model outputs to business actions.

4.3 DIMENSION TABLES

The provided content lists dimension tables for applicant details, credit history, device information, and time. These dimensions help the system analyze patterns by borrower segment, credit background, digital behavior, and temporal trends.

4.4 ADVANTAGES OF WAREHOUSE DESIGN

The user-provided content highlights three main advantages of the warehouse design: faster queries, better analytics, and scalability. In practical terms, the warehouse also improves reporting consistency and helps connect predictive outputs to downstream BI and dashboard use cases.

5. DATASET DESCRIPTION

The project relies on a synthetic dataset designed to simulate real-world lending behavior, and the attached content states that it contains more than 25,000 records. The presentation also groups features into affordability, stability, credit behavior, and digital risk signals, which provides a strong conceptual basis for the dataset chapter.

5.1 DATA ATTRIBUTES

The user-provided content lists core attributes such as Age, Income, Loan Amount, EMI, Employment Type, Credit History, Active Loans, and Delinquency. The presentation further expands feature groups to include affordability measures, stability indicators, credit behavior, and optional device or KYC risk flags.

5.2 DATA SIZE

The attached material specifies a dataset size of more than 25,000 records. This volume is sufficient for demonstrating preprocessing, model training, risk segmentation, and dashboard reporting in a student-level machine learning system.

5.3 DATA GENERATION

The project uses synthetic data to simulate lending behavior, and the attached content explicitly notes that the dataset is generated to reflect real-world credit and borrower patterns. Synthetic generation is useful in academic work because real credit data is highly sensitive and difficult to access due to privacy and regulatory constraints.

5.4 FEATURE GROUPS

The presentation organizes features into affordability, stability, credit behavior, and digital risk flags. It specifically mentions monthly income versus loan amount, EMI-to-income ratio, loan purpose, job tenure, employer type, residential stability, score band where available, active loan count, delinquency signals, new-device indicators, and KYC mismatch signals, with the note that only ethical and permitted signals are used.

5.5 DATA DICTIONARY NOTES

In the final report, this section should be expanded into a complete data dictionary table listing every feature name, description, data type, transformation applied, and whether it is used directly, derived through feature engineering, or only shown in reporting. This is an effective way to lengthen the report to 50 pages while improving academic quality.

6. SYSTEM REQUIREMENTS

The sample report includes a dedicated chapter for hardware and software requirements before moving deeper into technical implementation, and the same pattern is followed here. Since the present project is Python- and dashboard-based rather than cloud-streaming-based, the requirements are adapted accordingly.

6.1 HARDWARE REQUIREMENTS

A minimum Intel i5 or equivalent processor, 8 GB RAM, and 256 GB storage are appropriate for model experimentation, notebook execution, and dashboard testing, following the same

practical student development baseline seen in the sample report's hardware chapter. For smoother model training and visualization, 16 GB RAM is recommended.

6.2 SOFTWARE REQUIREMENTS

The software stack includes Python, Jupyter Notebook or Visual Studio Code, pandas, numpy, scikit-learn, XGBoost, SHAP, Streamlit, and standard plotting libraries. This requirement profile is consistent with the attached presentation's emphasis on ML, explainability, and interactive dashboard presentation.

7. LITERATURE SURVEY

The attached presentation contrasts traditional scorecards with ML-based scoring, highlights thin-file borrower challenges, and emphasizes explainability and monitoring, which together form the basis for a compact literature survey chapter. The sample report also reserves a complete chapter for literature before methodology, and the same reporting convention is retained. +1

7.1 TRADITIONAL CREDIT SCORING

The presentation notes that scorecards and Logistic Regression with WOE and IV remain industry standards because they are interpretable and auditable. Traditional scorecards are especially useful in regulated lending because they provide a transparent basis for explaining why an applicant was evaluated as higher or lower risk.

7.2 ML-BASED CREDIT RISK MODELS

The presentation states that Random Forest, XGBoost, and LightGBM consistently outperform traditional methods on tabular credit datasets. These algorithms are valuable because they can learn non-linear feature interactions and complex borrower patterns that simpler linear models may fail to capture.

7.3 THIN-FILE BORROWER CHALLENGES

The attached material emphasizes that millions of Indians are new-to-credit and may lack strong bureau history, which limits the usefulness of conventional underwriting alone. This challenge motivates the inclusion of alternative but ethical behavioral, stability, and device-linked indicators in risk modeling.

7.4 EXPLAINABILITY IN LENDING

The presentation highlights SHAP and LIME-style explanation approaches as important for decision transparency and model trust. In regulated or customer-sensitive environments, explanation is not an optional visualization feature; it is part of the system's accountability layer.

7.5 RESEARCH GAP AND NOVELTY

The attached content says the project focuses on PR-AUC, Recall, and KS statistic rather than relying only on accuracy, uses time-based splits to reduce leakage, and emphasizes business-aware threshold tuning and SHAP-driven explanations aligned with digital lending expectations. These elements form the project's novelty because they make the system more deployment-aware and operationally realistic than a basic classroom classification exercise.

8. TECHNOLOGY STACK

The project uses a fintech-oriented analytical stack that combines data engineering, machine learning, explainability, and lightweight deployment. This mirrors the chapterized technology discussion used in the sample report.

8.1 PYTHON

Python serves as the primary programming language for data preprocessing, model training, evaluation, and dashboard logic. It is especially suitable for credit analytics because of its extensive ecosystem for tabular learning and rapid experimentation.

8.2 DATA ANALYSIS LIBRARIES

Pandas and numpy support dataset manipulation, feature transformation, missing-value handling, and numerical analysis. These libraries form the foundation of the data processing workflow.

8.3 MACHINE LEARNING LIBRARIES

Scikit-learn supports baseline modeling, preprocessing, metrics, and evaluation utilities, while the attached presentation positions XGBoost and LightGBM as advanced model candidates for higher predictive performance.

8.4 EXPLAINABILITY TOOLS

The presentation explicitly identifies SHAP for both global feature importance and individual explanation delivery. SHAP helps connect abstract model outputs to user-understandable reason codes and business interpretation.

8.5 STREAMLIT

The user-provided content lists Streamlit as the dashboard layer for real-time prediction, risk visualization, and explanation display. Streamlit is appropriate for academic prototypes because it allows rapid development of interactive interfaces without heavy front-end engineering.

8.6 INTEGRATION OF TECHNOLOGIES

The overall stack integrates data processing, predictive modeling, explanation generation, decision logic, and dashboard display into one system. This complete workflow supports both academic demonstration and realistic fintech simulation.

9. METHODOLOGY

The attached presentation includes a proposed methodology pipeline covering data cleaning, feature engineering, baseline modeling, advanced modeling, and imbalance handling, while

also emphasizing leakage control through time-based splits. This chapter expands that pipeline into a report-ready methodology sequence.

9.1 DATA CLEANING

The user-provided content explicitly lists handling missing values and removing inconsistencies as part of preprocessing. Cleaning ensures that the dataset is reliable enough for training and prevents the model from learning from corrupted or misleading records.

9.2 ENCODING

The attached content states that categorical variables are converted into numerical form during preprocessing. Depending on the final implementation, this may involve label encoding, one-hot encoding, or WOE-based transformation for scorecard-style baseline models.

9.3 FEATURE ENGINEERING

The project specifically mentions feature engineering steps such as EMI-to-income ratio and broader risk indicators. The presentation also groups features into affordability, stability, credit behavior, and digital signals, which suggests a business-aware approach to engineered variable design.

9.4 TRAIN-TEST STRATEGY

The presentation states that the project uses time-based train-test splits to avoid leakage and obtain more realistic out-of-time performance estimates. This is a strong methodological choice because random splits can artificially inflate model quality in financial datasets.

9.5 IMBALANCE HANDLING

The methodology pipeline in the presentation explicitly includes imbalance handling. This is important because default datasets are usually skewed, and evaluation must account for the fact that high accuracy can be misleading when bad cases are relatively rare.

9.6 BASELINE MODEL

The presentation identifies Logistic Regression scorecards with WOE and IV as the transparent, auditable benchmark model. A baseline of this type is useful because it establishes a regulatory-friendly reference point before introducing more complex learners.

9.7 ADVANCED MODEL

The attached content states that XGBoost and LightGBM are used as advanced models to capture non-linear interactions and outperform the baseline on PR-AUC. These models are suitable when the data contains complex relationships among affordability, stability, and behavioral variables.

9.8 EXPLAINABILITY WORKFLOW

The methodology includes generation of SHAP-based global feature importance and per-user explanation outputs, as directly described in the presentation. This ensures that predictions are accompanied by reasons rather than appearing as opaque scores.

9.9 END-TO-END PIPELINE

The proposed methodology pipeline in the attached presentation is described as a clean, auditable sequence from raw data to explainable credit decisions. In report form, the complete flow can be summarized as follows: collect or generate synthetic applicant data, clean and encode it, engineer business-relevant features, split the data temporally, train baseline and advanced models, evaluate them with suitable credit metrics, apply threshold tuning, assign risk bands, and present the final decision with SHAP-based explanation.

10. MACHINE LEARNING MODELS

The user-provided content identifies Logistic Regression as the baseline model and XGBoost as the main advanced model, while the presentation also mentions LightGBM as another strong candidate. Together, these models create a useful balance between interpretability and predictive power.

10.1 LOGISTIC REGRESSION

The attached content describes Logistic Regression as a baseline model and notes its interpretability. It is valuable in credit systems because the model's linear structure makes relationships easier to explain, audit, and benchmark.

10.2 XGBOOST

The attached content describes XGBoost as an advanced model with high accuracy and strong ability to handle non-linearity. The presentation also suggests that XGBoost is expected to outperform the baseline on PR-AUC through better generalization on complex tabular data.

10.3 LIGHTGBM

The presentation mentions LightGBM as part of the advanced modeling stage alongside XGBoost. It provides another gradient boosting approach that is well suited to structured lending data and efficient experimentation.

10.4 MODEL COMPARISON LOGIC

The evaluation plan shown in the presentation compares Logistic Regression, XGBoost, and LightGBM while emphasizing credit-appropriate metrics such as PR-AUC, Recall, KS statistic, and threshold tuning. This comparative design is stronger than using a single model because it lets the project justify why one learner is preferred in the final engine.

11. RISK BANDING AND DECISION ENGINE

The user-provided content includes a clear probability-to-risk mapping: PD below 0.30 is Low risk and Approve, 0.30 to 0.60 is Medium risk and Review, and above 0.60 is High risk and Reject. This chapter formalizes that mapping as the business layer of the system.

11.1 PD THRESHOLD DESIGN

The system converts raw model output into risk bands using predefined probability thresholds. This is important because business teams make decisions more easily from interpretable categories than from raw probabilities alone.

11.2 LOW, MEDIUM, HIGH RISK MAPPING

The attached content defines the mapping as follows: probability of default below 0.30 corresponds to Low risk and Approve, values from 0.30 to 0.60 correspond to Medium risk and Review, and values above 0.60 correspond to High risk and Reject. These thresholds can be tuned later based on portfolio loss tolerance and approval strategy.

PD Range	Risk Band	Decision
< 0.30	Low	Approve
0.30–0.60	Medium	Manual Review
> 0.60	High	Reject

11.3 APPROVAL LOGIC

The attached content says the system converts model output into actionable decisions such as Approve, Manual Review, and Reject. This decision engine is the point where analytics becomes operational policy.

11.4 BUSINESS INTERPRETATION

The presentation emphasizes business-aware threshold tuning because optimal cut-offs should balance default capture against rejection volume and expected loss. This means threshold selection should reflect lending economics rather than generic classification performance alone.

12. EXPLAINABILITY (SHAP)

The attached content identifies explainability as a core system requirement and specifically lists SHAP under project scope and deliverables. This makes explainability a first-class module rather than an optional appendix.

12.1 PURPOSE

The user-provided content states that explainability is used to understand model decisions and improve transparency. In credit systems, this is important for internal trust, audit readiness, and customer-facing communication.

12.2 GLOBAL EXPLANATIONS

The attached presentation explicitly mentions global feature importance as one of the explainability outputs. Global explanation helps analysts understand which variables consistently influence default prediction across the portfolio.

12.3 INDIVIDUAL REASON CODES

The presentation states that the project delivers per-user SHAP reason codes and individual explanations. This is especially useful when the system needs to justify why a specific applicant received a review or rejection outcome.

12.4 COMPLIANCE RELEVANCE

The presentation connects explainable decisions with compliance and digital lending expectations, stating that SHAP-driven explanations make outputs actionable and aligned with RBI-oriented lending transparency needs. This makes the explainability layer both technically and operationally valuable.

13. DASHBOARD DESIGN

The user-provided content describes a Streamlit dashboard with user input, real-time prediction, risk visualization, and an explanation panel. This section expands that interface into a report-ready design narrative.

13.1 STREAMLIT INTERFACE

The attached content identifies Streamlit as the main dashboard framework. It is suitable for this project because it supports direct integration with Python-based predictive logic and quick rendering of decision outputs.

13.2 INPUT MODULE

The dashboard contains a user input form where loan details and applicant attributes are entered, as described in the workflow section of the attached content. Typical fields include age, income, EMI burden, employment type, loan request, and selected behavioral signals.

13.3 PREDICTION OUTPUT MODULE

The user-provided content lists PD Score, Risk Level, and Decision as core outputs. These outputs allow the dashboard to function as a compact decision-support interface for lending analysis.

13.4 VISUALIZATION AND EXPLANATION PANEL

The attached content states that the dashboard provides risk visualization and explanation display, and the UI is described as a light-themed, card-based layout with clean spacing and dynamic updates. The dashboard can therefore be documented as both functionally interactive and visually structured for readability.

13.5 REPORT EXPORT OPTIONS

The user-provided content mentions CSV export and PDF report generation as output modules. These export options are useful for audit trails, stakeholder review, and case-level documentation.

14. IMPLEMENTATION

The implementation chapter should present the system as a sequence of practical development tasks, similar to how the sample report describes environment setup, module implementation, integration, and testing. The actual code details can be inserted later as screenshots, notebook snippets, or module screenshots.

14.1 ENVIRONMENT SETUP

The implementation begins with Python environment setup and installation of the required libraries for data handling, modeling, explainability, and dashboard execution. Jupyter Notebook or VS Code may be used during development, while Streamlit is used to run the final interface.

14.2 DATA PREPARATION

Data preparation includes loading the synthetic dataset, checking schema consistency, cleaning missing or inconsistent values, encoding categorical variables, and creating engineered risk indicators such as EMI-to-income ratio. This stage directly reflects the preprocessing steps listed in the user-provided content.

14.3 MODEL TRAINING

The next stage trains the baseline and advanced models using the prepared feature set. The presentation emphasizes Logistic Regression as a transparent benchmark and XGBoost or LightGBM as advanced learners for stronger predictive performance.

14.4 THRESHOLD TUNING

The presentation highlights threshold tuning as part of the evaluation logic because credit decisions must balance missed bad loans against excessive rejections. This tuning stage converts raw model probabilities into operational risk bands and decision outcomes.

14.5 DASHBOARD INTEGRATION

The trained model, threshold logic, and explanation layer are integrated into the Streamlit dashboard so that a user can enter inputs and immediately receive predicted risk information. This completes the transition from offline model experimentation to interactive decision support.

14.6 TESTING AND VALIDATION

Testing includes checking input handling, scoring consistency, threshold mapping, explanation visibility, and export behavior. Additional validation can include comparison of candidate models and repeated testing on out-of-time samples, consistent with the presentation's focus on realistic evaluation.

15. MODEL EVALUATION

The attached presentation and the user-provided content emphasize evaluation metrics such as Accuracy, Recall, Precision, PR-AUC, and KS Statistic, with a clear statement that XGBoost performs better because it handles complex patterns and generalizes more effectively. This chapter organizes those points into formal report analysis.

15.1 EVALUATION METRICS

The presentation identifies PR-AUC as especially important for imbalanced default datasets, Recall as important for capturing true default cases, KS Statistic as an industry standard for score separation, and threshold tuning as a business-aware control point. It also links model output to expected loss concepts such as PD, LGD, and EAD.

15.2 WHY PR-AUC AND KS MATTER

The attached presentation explicitly notes that standard accuracy can be misleading on imbalanced default data and that the project therefore focuses on PR-AUC, Recall, and KS Statistic. This is a strong credit-risk evaluation choice because it rewards a model for identifying rare bad cases rather than simply labeling most cases as safe.

15.3 EXPECTED RESULTS

The presentation lists expected deliverables such as PD risk score, Low-Medium-High risk bands, better PR-AUC than baseline, top feature drivers, SHAP-based top reasons, a clean dataset, a data dictionary, and an evaluation report. These expected outputs can be turned into tables, screenshots, and appendix material to help the report reach a full 50-page format.

15.4 REAL-WORLD LENDING INTERPRETATION

The presentation includes a real-world example of a UPI-first fintech lending scenario in which a sample applicant receives a predicted PD of 0.23, a Medium Risk label, and an action such as approval with lower limit or manual review, supported by SHAP-style top reasons. This example demonstrates how the engine can translate model output into business actions rather than leaving risk scores uninterpreted.

16. APPLICATIONS

The user-provided content states that the project can be applied to digital lending platforms, NBFC systems, and credit scoring tools. The presentation also mentions BNPL platforms, salary advance applications, embedded finance, and partner-bank co-lending programs as real-world use cases.

These applications show that FINOVA is not limited to one product format. Its architecture is flexible enough to support different short-cycle lending contexts where fast decisions and explainability are both important.

17. ADVANTAGES

The user-provided content lists automated decision-making, high accuracy, scalability, and explainability as key advantages. The presentation further supports these benefits by emphasizing deployment readiness, auditable benchmarking, strong tabular performance, and SHAP-based explanation.

Additional advantages include business-oriented risk banding, support for thin-file users, and a dashboard layer that makes model output easier to interpret in operational settings. These strengths make the project academically solid and practically relevant.

18. LIMITATIONS

The user-provided content identifies three main limitations: synthetic dataset dependency, model assumptions, and the absence of a real-time API. These are reasonable limitations for an academic prototype because access to live lending data and production infrastructure is usually restricted.

Another practical limitation is that model performance in a synthetic environment may not fully represent real portfolio behavior under changing macroeconomic conditions. This is why the project should be framed as a realistic simulation and foundation rather than a fully production-validated credit policy engine.

19. FUTURE ENHANCEMENTS

The user-provided content lists several future enhancements, including real-time API integration, cloud deployment, advanced visualization, user authentication, and use of real datasets. The presentation's monitoring and deployment-oriented framing also supports expanding the system toward model monitoring, drift detection, and continuous policy calibration.

Additional enhancements can include bureau integration, document verification modules, dynamic pricing, fairness monitoring, automated retraining, and portfolio-level reporting

dashboards. These points can be expanded significantly in the final report to increase length and academic depth.

20. CONCLUSION

The user-provided content states that FINOVA demonstrates an end-to-end ML pipeline, data warehousing integration, and a realistic fintech simulation, while providing a strong foundation for production-level credit risk systems. That conclusion is consistent with the attached presentation, which frames the project as accurate, explainable, and practical for India's digital lending ecosystem.

The report structure also mirrors the broad chapter progression used in the uploaded sample project report, which makes the document suitable for academic expansion into a full submission-ready report with screenshots, figures, diagrams, model tables, and appendices. By combining warehouse design, ML scoring, explainability, and dashboard output, FINOVA stands as a strong applied analytics project in the fintech domain.

21. REFERENCES

The user-provided content lists the project's reference direction as machine learning algorithms, fintech systems, XGBoost documentation, and SHAP documentation. The attached presentation also references Logistic Regression with WOE and IV, XGBoost, LightGBM, SHAP, and lending-oriented evaluation concepts such as PR-AUC and KS Statistic.

Suggested references for the final version:

- XGBoost Documentation.
- SHAP Documentation.
- Scikit-learn Documentation.
- Literature on credit scoring and scorecard development.
- Research on digital lending and thin-file borrower modeling.
- RBI and digital lending guidance where relevant.